

Operating Systems

CSCI 3150

Lecture 2: Architectural (hardware) Support for OS

Hong Xu

Content of this lecture

- Review of introduction
- Hardware overview
- A peek at Unix
- Hardware (architecture) support
- Summary

Review

- What are the two main responsibilities of OS?
 - Manage **hardware** resources
 - Provide a clean set of interface to **programs**
- Managing resources:
 - Allocation
 - Protection
 - Reclamation
 - Virtualization
- Questions?

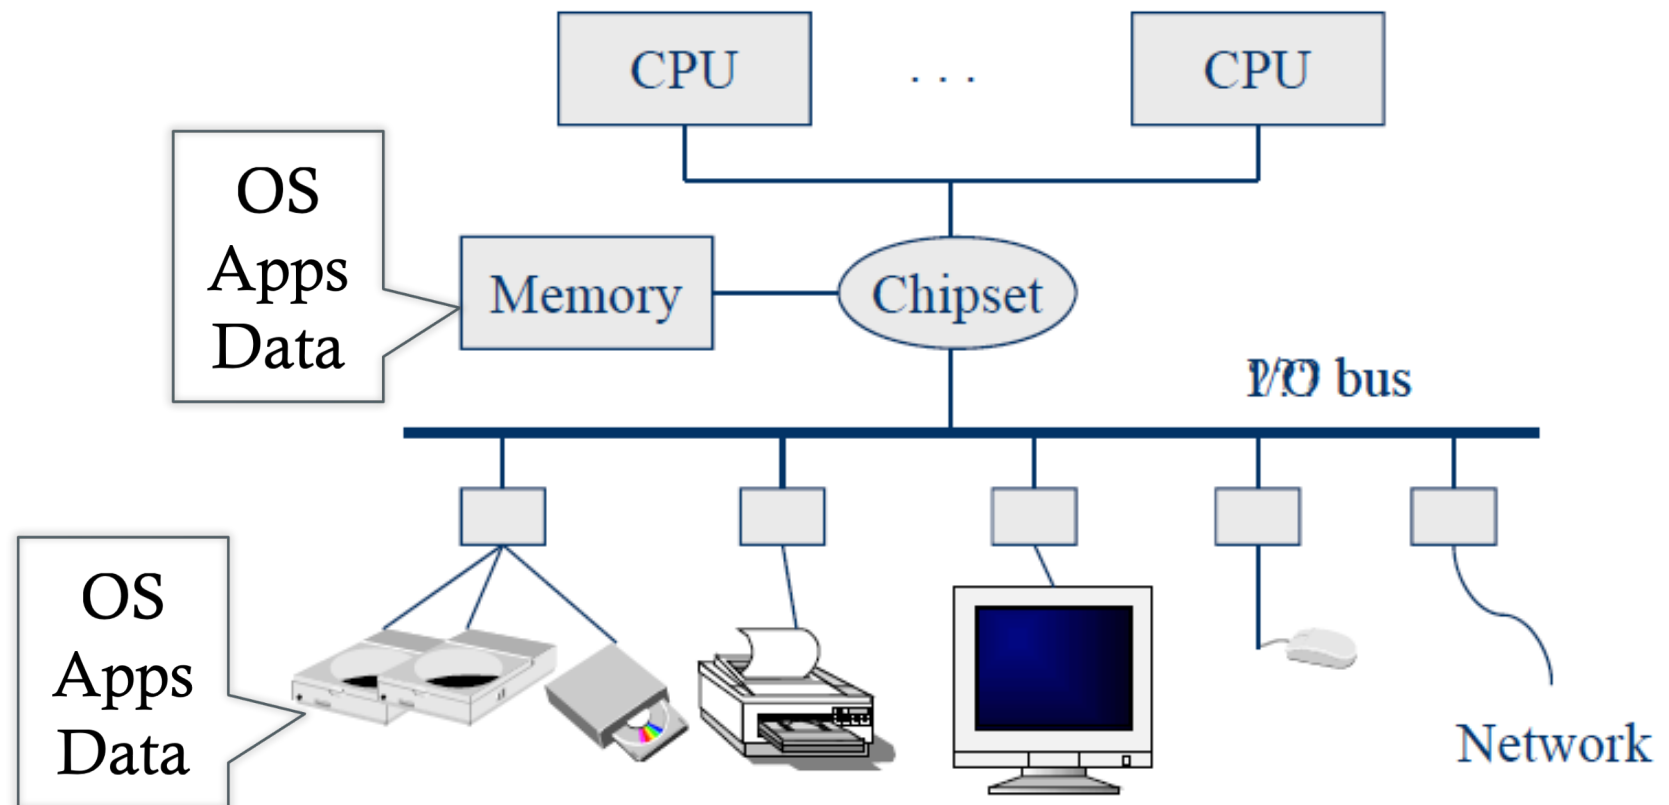
Why Start With Hardware?

- Operating system functionality fundamentally depends upon hardware
 - Key goal of an OS is to manage hardware
 - If done well, applications can be oblivious to HW details
- Hardware support can greatly simplify – or **complicate** – OS tasks
 - Early PC operating systems (DOS, MacOS) lacked virtual memory in part because the hardware did not support it
 - <https://github.com/microsoft/ms-dos>


So what is inside a computer

- An abstract overview
 - <http://www.youtube.com/watch?v=Q2hmuqS8bwM>
- An introduction with a real computer
 - <https://www.youtube.com/watch?v=HB4I2CgkcCo>

A Typical Computer from a Hardware Point of View



Memory-storage Hierarchy

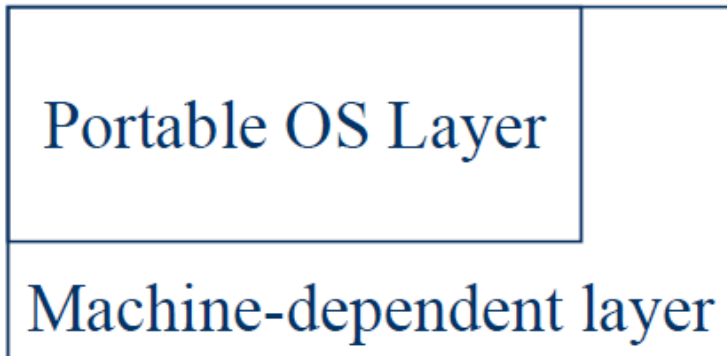


1 nanosecond = 10^{-9} second

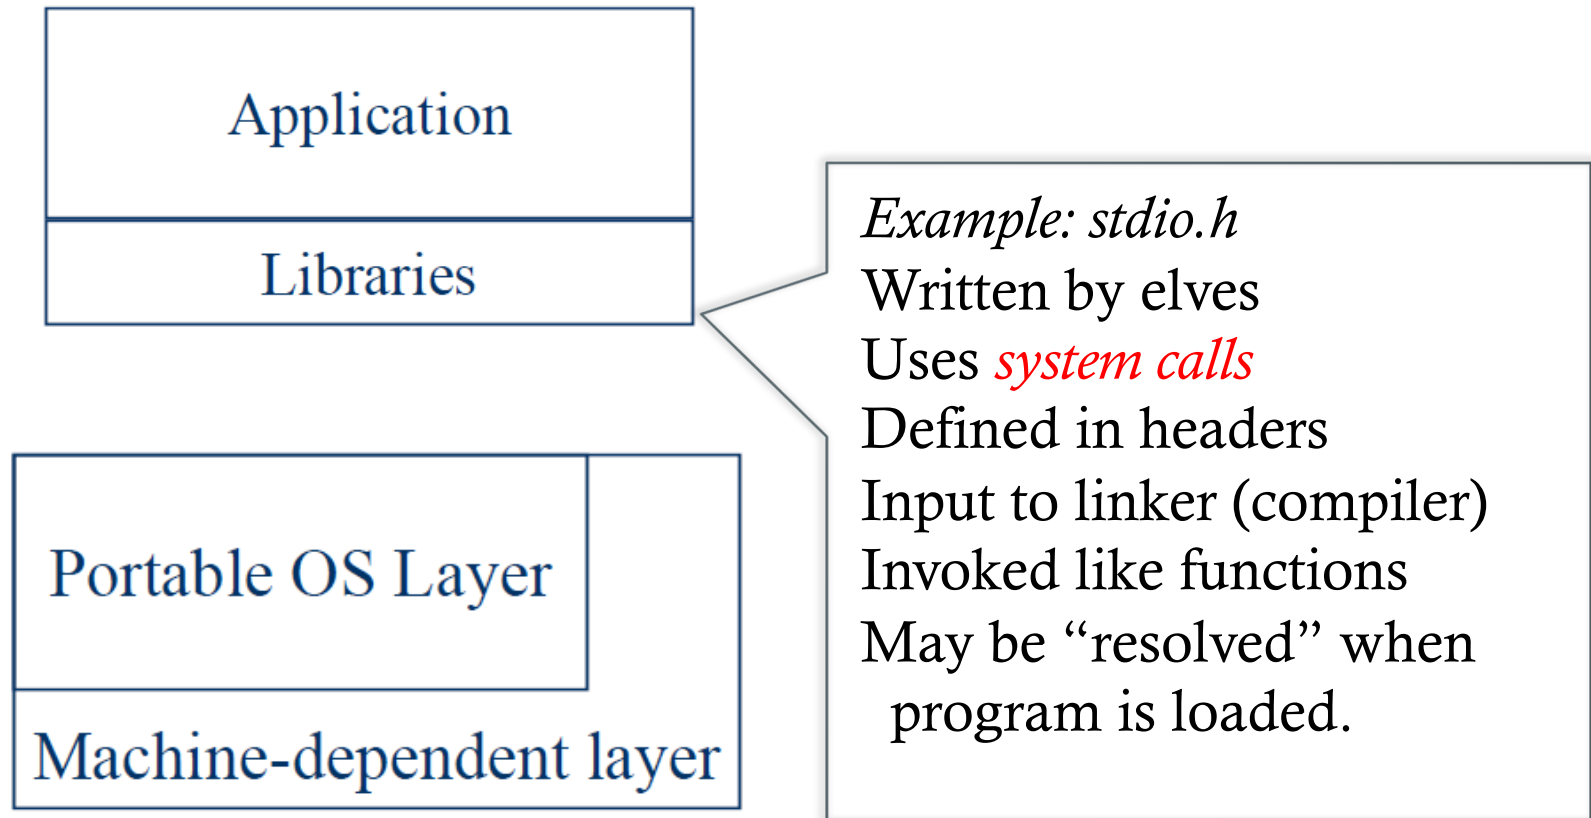
A peek into Unix structure



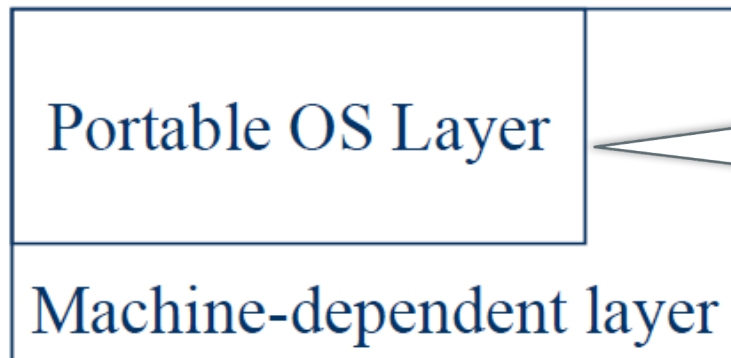
Written by programmer
Compiled by programmer
Uses library calls (e.g., printf)



A peek into Unix structure

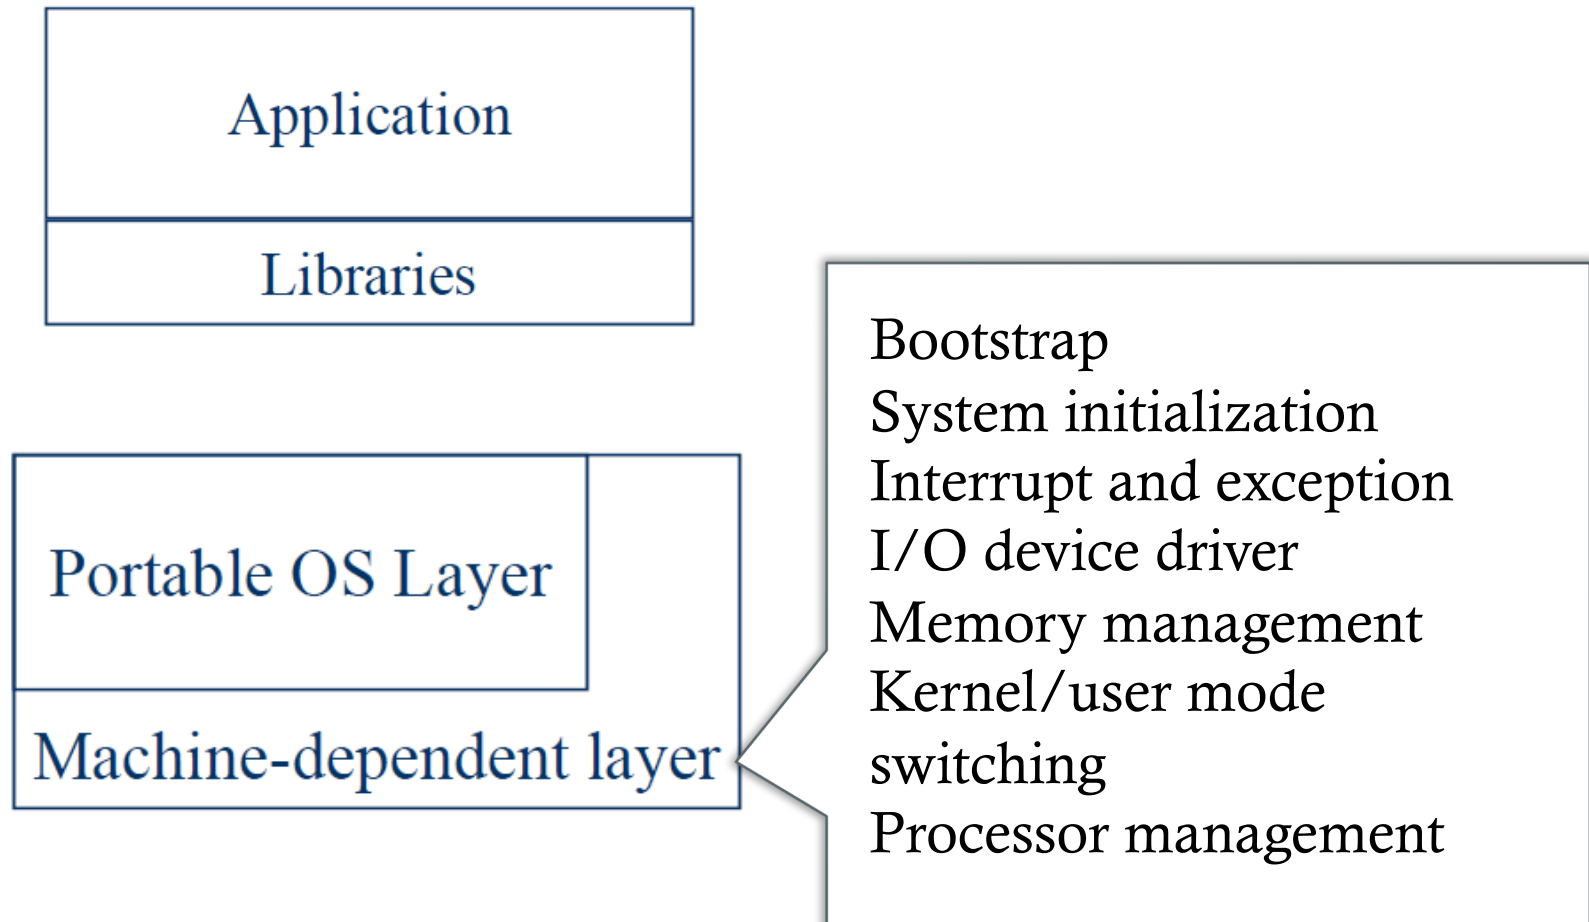


A peek into Unix structure



System calls (read, open..)
All “high-level” code

A peek into Unix structure

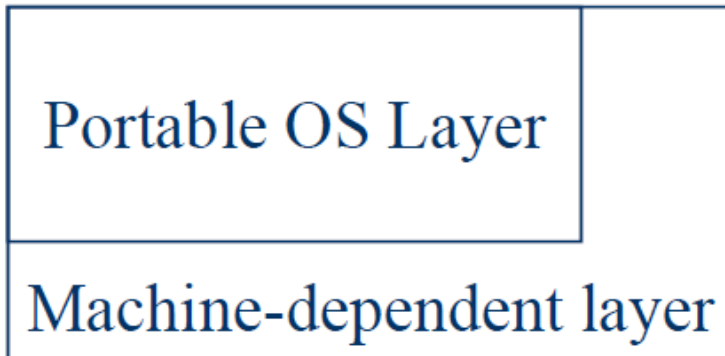


A peek into Unix structure



Cannot execute
“protected_instruction”, e.g.,
directly access I/O device

User mode



Kernel mode

- Some systems do not have clear user-kernel boundary
- User/kernel mode is supported by hardware (why?)

Why hardware has to support User/Kernel mode?

Imaginary OS code (software-only solution)

```
if ([PC] != protected_instruction)
    execute(PC);
else
    switch_to_kernel_mode();
```



Why hardware has to support User/Kernel mode?

Application's code:

```
lw    $t0, 4($gp)
mult  $t0, $t0, $t0
lw    $t1, 4($gp)
ori   $t2, $zero, 3
mult  $t1, $t1, $t2
add   $t2, $t0, $t1
sw    $t2, 0($gp)
```

OS: check if next instruction is protected instruction.

Why hardware has to support User/Kernel mode?

Application's code:

```
lw      $t0, 4($gp)
mult    $t0, $t0, $t0
lw      $t1, 4($gp)
ori      $t2, $zero, 3
mult     $t1, $t1, $t2
add      $t2, $t0, $t1
sw      $t2, 0($gp)
```

OS: check if next instruction is protected instruction.

- Performance overhead is too big: OS needs to check every instruction of the application!
- *Simulators*

Why hardware has to support User/Kernel mode?

OS: set-up the environment;
load the application

Application's code:

```
lw    $t0, 4($gp)
mult  $t0, $t0, $t0
lw    $t1, 4($gp)
ori   $t2, $zero, 3
mult  $t1, $t1, $t2
add   $t2, $t0, $t1
sw    $t2, 0($gp)
```

• Instead, what we really want is to give the CPU entirely to the application

• *Bare-metal execution*

- *Any problems?*
- *How can OS check if application executes protected instruction?*
 - *How can OS know it will ever run again?*

Return to OS after termination;
OS: schedule next application to execute..

Why hardware has to support User / Kernel mode?

- Give the CPU to the user application
 - **Why: Performance and efficiency**
 - OS will not be executing
- Without hardware's help, OS loses control of the machine!
 - *Analogy: give the car key to someone, how do you know if he will return the car?*
- *This is the most fundamental reason why OS will need hardware support --- not only for user / kernel mode*

Questions?

Hardware Features for OS

- Features that directly support the OS include
 - Protection (kernel/user mode)
 - Protected instructions
 - Memory protection
 - System calls
 - Interrupts and exceptions
 - Timer (clock)
 - I/O control and operation
 - Synchronization

Types of Hardware Support

- Manipulating privileged machine state
 - Protected instructions
 - Manipulate device registers, TLB entries, etc.
- Generating and handling “events”
 - Interrupts, exceptions, system calls, etc.
 - Respond to external events
 - CPU requires software intervention to handle fault or trap
- Mechanisms to handle concurrency
 - Interrupts, atomic instructions

Protected Instructions

- A subset of instructions of every CPU is restricted to use only by the OS
 - Known also as privileged instructions
- Only the operating system can
 - Directly access I/O devices (disks, printers, etc.)
 - Security, fairness (why?)
 - Manipulate memory management state
 - Page table pointers, page protection, TLB management, etc.
 - Manipulate protected control registers
 - Kernel mode, interrupt level
 - Halt instruction (why?)

OS Protection

- Hardware must support (at least) two modes of operation: **kernel** mode and **user** mode
 - Mode is indicated by a status bit in a protected control register
 - User programs execute in user mode
 - OS executes in kernel mode (OS == “kernel”)
- Protected instructions only execute in kernel mode
 - **CPU** checks mode bit when protected instruction executes
 - Setting mode bit must be a protected instruction
 - Attempts to execute in user mode are detected and prevented
 - x86: General Protection Fault


Memory Protection

- OS must be able to protect programs from each other
- OS must protect itself from user programs
- We need hardware support
 - *Again: once OS gives the CPU to the user programs, OS loses control*

Memory Protection

- Memory management hardware provides memory protection mechanisms
 - Base and limit registers
 - Page table pointers, page protection, TLB
 - Virtual memory
 - Segmentation
- Manipulating memory management hardware uses privileged operations

Hardware Features for OS

- Features that directly support the OS include
 - Protection (kernel/user mode)
 - Privileged instructions
 - Memory protection
 - System calls
 - Interrupts and exceptions
 - Timer (clock)
 - I/O control and operation
 - Synchronization
- 
- Questions?**

Events

- After the OS has booted, **all entry to the kernel happens as the result of an event**
 - event immediately stops current execution
 - changes mode to kernel mode, event handler is called
- An event is an “unnatural” change in control flow
 - Events immediately stop current execution
 - Changes mode, context (machine state), or both
- The kernel defines a handler for each event type
 - Event handlers always execute in kernel mode
 - The specific types of events are defined by the machine
- **In effect, the operating system is one big event handler**

OS Control Flow

- When the processor receives an event of a given type, it
 - transfers control to handler within the OS
 - handler saves program state (PC, registers, etc.)
 - handler functionality is invoked
 - handler restores program state, returns to program

Categorizing Events

- Two kinds of events, **interrupts** and **exceptions**
- Exceptions are caused by executing instructions
 - CPU requires software intervention to handle a fault or trap
- Interrupts are caused by an external event
 - Device finishes I/O, timer expires, etc.
- Two *reasons* for events, **unexpected** and **deliberate**
- Unexpected events are, well, unexpected
 - **What is an example?**
- Deliberate events are scheduled by OS or application
 - **Why would this be useful?**

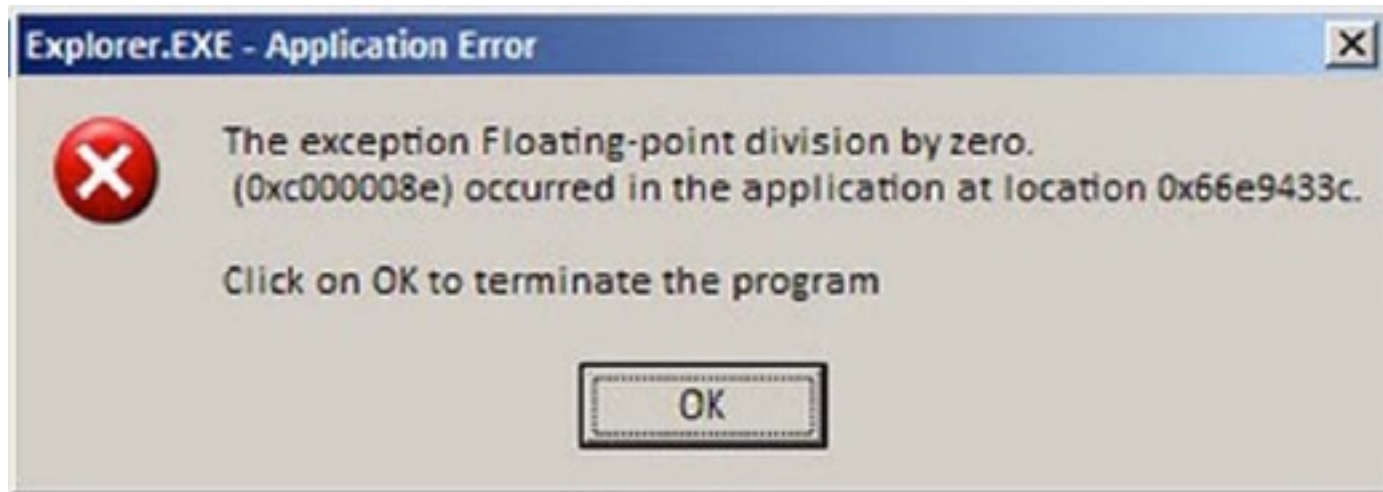
Categorizing Events

- This gives us a convenient table:

	Unexpected	Deliberate
Exceptions (sync)	fault	syscall trap
Interrupts (async)	interrupt	software interrupt

- Terms may be used slightly differently by various OSes, CPU architectures...
 - No need to “memorize” all the terms
- Software interrupt – a.k.a. async system trap (AST), async or deferred procedure call (APC or DPC)
- Will cover faults, system calls, and interrupts next

Faults



Faults

- Hardware detects and reports “exceptional” conditions
 - Page fault, divide by zero, unaligned access
- Upon exception, hardware “faults” (verb)
 - Must save state (PC, registers, mode, etc.) so that the faulting process can be restarted
- Fault exceptions are a performance optimization
 - Could detect faults by inserting extra instructions into code (at a significant performance penalty)

Handling Faults

- Some faults are handled by “fixing” the exceptional condition and returning to the faulting context
 - Page faults cause the OS to place the missing page into memory
 - Fault handler resets PC of faulting context to re-execute instruction that caused the page fault
- Some faults are handled by notifying the process
 - Fault handler changes the saved context to transfer control to a user-mode handler on return from fault
 - Handler must be registered with OS
 - Unix **signals**
 - SIGSEGV, SIGALRM, SIGTERM, etc.

Handling Faults

- The kernel may handle unrecoverable faults by killing the user process
 - Program faults with no registered handler
 - Halt process, write process state to file, destroy process
 - In Unix, the default action for many signals (e.g., SIGSEGV)
- What about faults in the kernel?
 - Dereference NULL, divide by zero, undefined instruction
 - These faults considered fatal, operating system crashes
 - **Unix panic**, **Windows “Blue screen of death”**
 - Kernel is halted, state dumped to a core file, machine locked up

System Calls

- For a user program to do something “privileged” (e.g., I/O) it must call an OS procedure
 - Known as **crossing the protection boundary**, or a **protected procedure call**
- Hardware provides a **system call** instruction that:
 - Causes an exception, which vectors to a kernel handler
 - Passes a parameter determining the system routine to call
 - Saves caller state (PC, registers, etc.) so it can be restored
 - Returning from system call restores this state
- Requires hardware support to:
 - Restore saved state, reset mode, resume execution

System Call Functions

- Process control
 - Create process, allocate memory
- File management
 - Create, read, delete file
- Device management
 - Open device, read/write device, mount device
- Information maintenance
 - Get time
- Programmers generally do not use system calls directly
 - They use runtime libraries (e.g., `stdio.h`)
 - *Why?*

Function call

```
main () {  
    foo (10);  
}
```

Compile



```
main:    push $10  
         call foo  
         .. ..  
foo:     .. ..  
         ret
```

System call

open (path, flags, mode);

More info:

<https://stackoverflow.com/questions/1817577/what-does-int-0x80-mean-in-assembly-code>

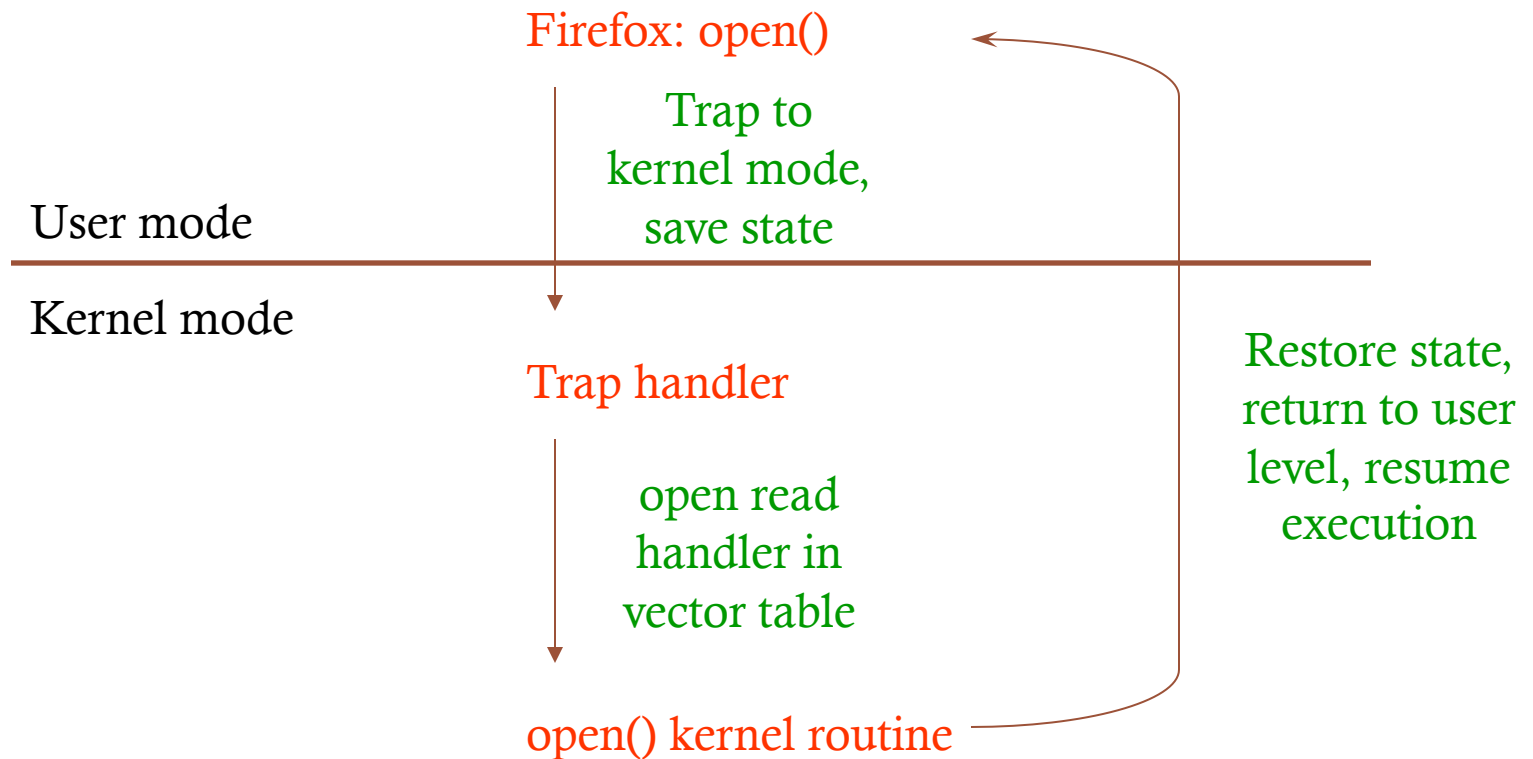
```
open: ;Linux convention:
      ;parameters via registers.
      mov eax, 5 ; syscall number for open
      mov ebx, path ; ebx: first parameter
      mov ecx, flags ; ecx: 2nd parameter
      mov edx, mode ; edx: 3rd parameter
      int 80h
```

```
open: ; FreeBSD convention:
      ; parameters via stacks.
      push dword mode
      push dword flags
      push dword path
      mov eax, 5
      push dword eax ; syscall number
      int 80h
      add esp, byte 16
```

Directly using system call?

- Write assembly code
 - Hard
- Poor portability
 - write different version for different architecture
 - write different version for different OSes
- Application programmers use library
 - Libraries written by elves


System Call



Steps in making a syscall

Example:

`read (fd, buffer, nbytes)`



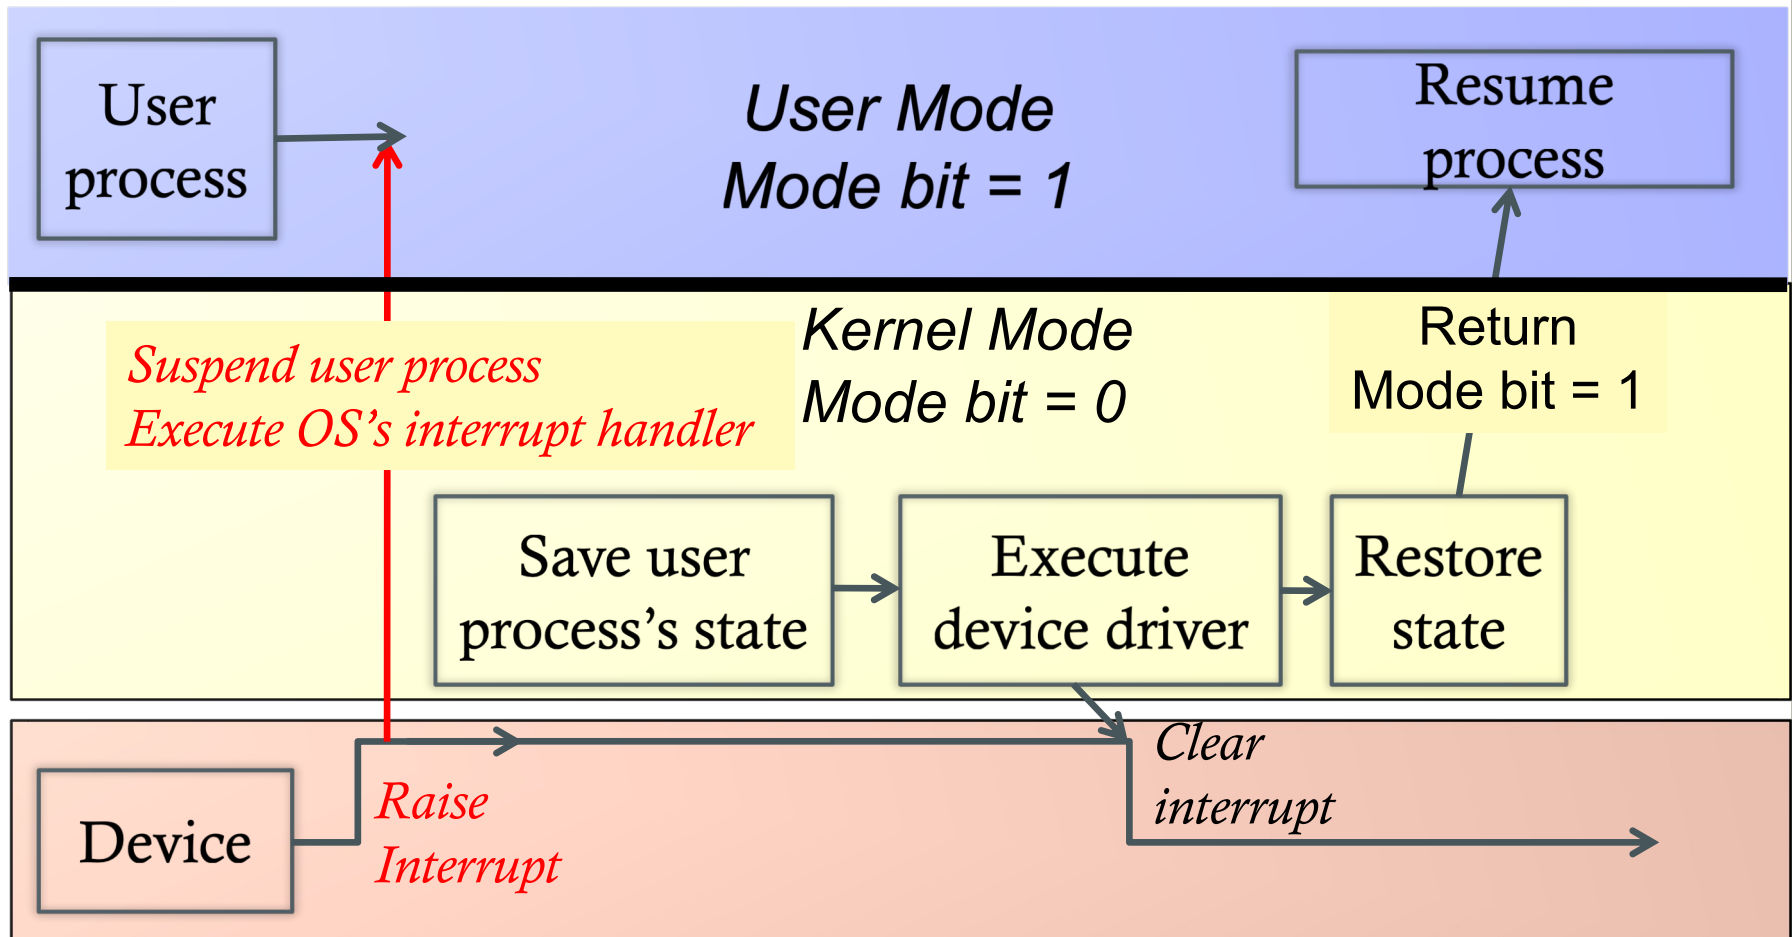
System Call Issues

- What would happen if the kernel did not save state?
- Why must the kernel verify arguments?
- Why is a table of system calls in the kernel necessary?

Interrupts

- Interrupts signal asynchronous events
 - I/O hardware interrupts
 - Hardware timers

Interrupt Illustrated



How to find interrupt handler?

- Hardware maps interrupt type to interrupt number
- OS sets up Interrupt Descriptor Table (IDT) at boot
 - Also called interrupt vector
 - IDT is in memory
 - Each entry is an interrupt handler
 - OS lets hardware know IDT base
- Hardware finds handler using interrupt number as index into IDT
 - `handler = IDT[intr_number]`

Timer

- The timer is critical for an operating system
- It is the fallback mechanism by which the OS reclaims control over the machine
 - Timer is set to generate an interrupt after a period of time
 - Setting timer is a privileged instruction
 - When timer expires, generates an interrupt
 - Handled by kernel, which controls resumption context
 - Basis for OS [scheduler](#) (*more later...*)
- Prevents infinite loops
 - OS can always regain control from erroneous or malicious programs that try to hog CPU
- Also used for time-based functions (e.g., *sleep()*)

I/O Control

- I/O issues
 - Initiating an I/O
 - Completing an I/O
- Initiating an I/O
 - Special instructions
 - Memory-mapped I/O
 - Device registers mapped into address space
 - Writing to address sends data to I/O device

I/O Completion

- Interrupts are the basis for asynchronous I/O
 - OS initiates I/O
 - Device operates independently of rest of machine
 - Device sends an interrupt signal to CPU when done
 - OS maintains a vector table containing a list of addresses of kernel routines to handle various events
 - CPU looks up kernel address indexed by interrupt number, context switches to routine

I/O Example

1. Ethernet receives packet, writes packet into memory
2. Ethernet signals an interrupt
3. CPU stops current operation, switches to kernel mode, saves machine state (PC, mode, etc.) on kernel stack
4. CPU reads address from vector table indexed by interrupt number, branches to address (Ethernet device driver)
5. Ethernet device driver processes packet (reads device registers to find packet in memory)
6. Upon completion, restores saved state from stack

Interrupt Questions

- Interrupts halt the execution of a process and transfer control (execution) to the operating system
 - Can the OS be interrupted? (Consider why there might be different IRQ levels)
- Interrupts are used by devices to notify the OS that something needs attention
 - What is an alternative approach to achieve this?
 - What are the drawbacks of that approach?

Alternative approach

- Polling

```
while (Ethernet_card_queue_is_empty)
    ;
// Ethernet card received packets.
handle_packets();
```

- Problems?

- Analogy:

- Polling: keeps checking the email every 30 seconds
- Interrupt: when email arrives, give me a ring

Summary

- Protection
 - User/kernel modes
 - Protected instructions
- System calls
 - Used by user-level processes to access OS functions
 - Access what is “in” the OS
- Exceptions
 - Unexpected event during execution (e.g., divide by zero)
- Interrupts
 - Timer, I/O

Summary (2)

- After the OS has booted, **all entry to the kernel happens as the result of an event**
 - event immediately stops current execution
 - changes mode to kernel mode, event handler is called
- When the processor receives an event of a given type, it
 - transfers control to handler within the OS
 - handler saves program state (PC, registers, etc.)
 - handler functionality is invoked
 - handler restores program state, returns to program

Architecture Trends Impact OS Design

- Processor
 - Single core to multi-core
 - OS must better handle concurrency
- Network
 - Isolation to dial-up to LAN to WAN
 - OS must devote more efforts to communications
 - Disconnected to wired to wireless
 - OS must manage connectivity more
 - Isolated to shared to attacked
 - OS must provide more security/protection
- Mobile/battery-operated
 - OS must pay attention to energy consumption

May you live in interesting times

- Multicores
- Smart phones
- Tapes → disks → flash memory → ..
- 3G, 4G, 5G.
- Cloud
- Wearable computers
- Virtual reality
- Motion capturing device
- ..