

# Operating Systems

## CSCI 3150

### *Lecture 4: Threads*

Hong Xu

# Processes

- Recall that a process includes many things
  - An address space (defining all the code and data pages)
  - OS resources (e.g., open files) and accounting information
  - Execution state (PC, SP, regs, etc.)
- Creating a new process is costly because of all of the data structures that must be allocated and initialized
  - Recall struct proc in Solaris
  - ...which does not even include page tables, perhaps TLB flushing, etc.
- Communicating between processes is costly because most communication goes through the OS
  - Overhead of system calls and copying data

# Parallel Programs

- To execute these programs we need to
  - Create several processes that execute in parallel
  - Cause each to map to the same address space to share data
    - They are all part of the same computation
  - Have the OS schedule these processes in parallel
- This situation is **very inefficient**
  - **Space**: PCB, page tables, etc.
  - **Time**: create data structures, fork and copy addr space, etc.
- Solutions: possible to have more **efficient**, yet **cooperative** “processes”?

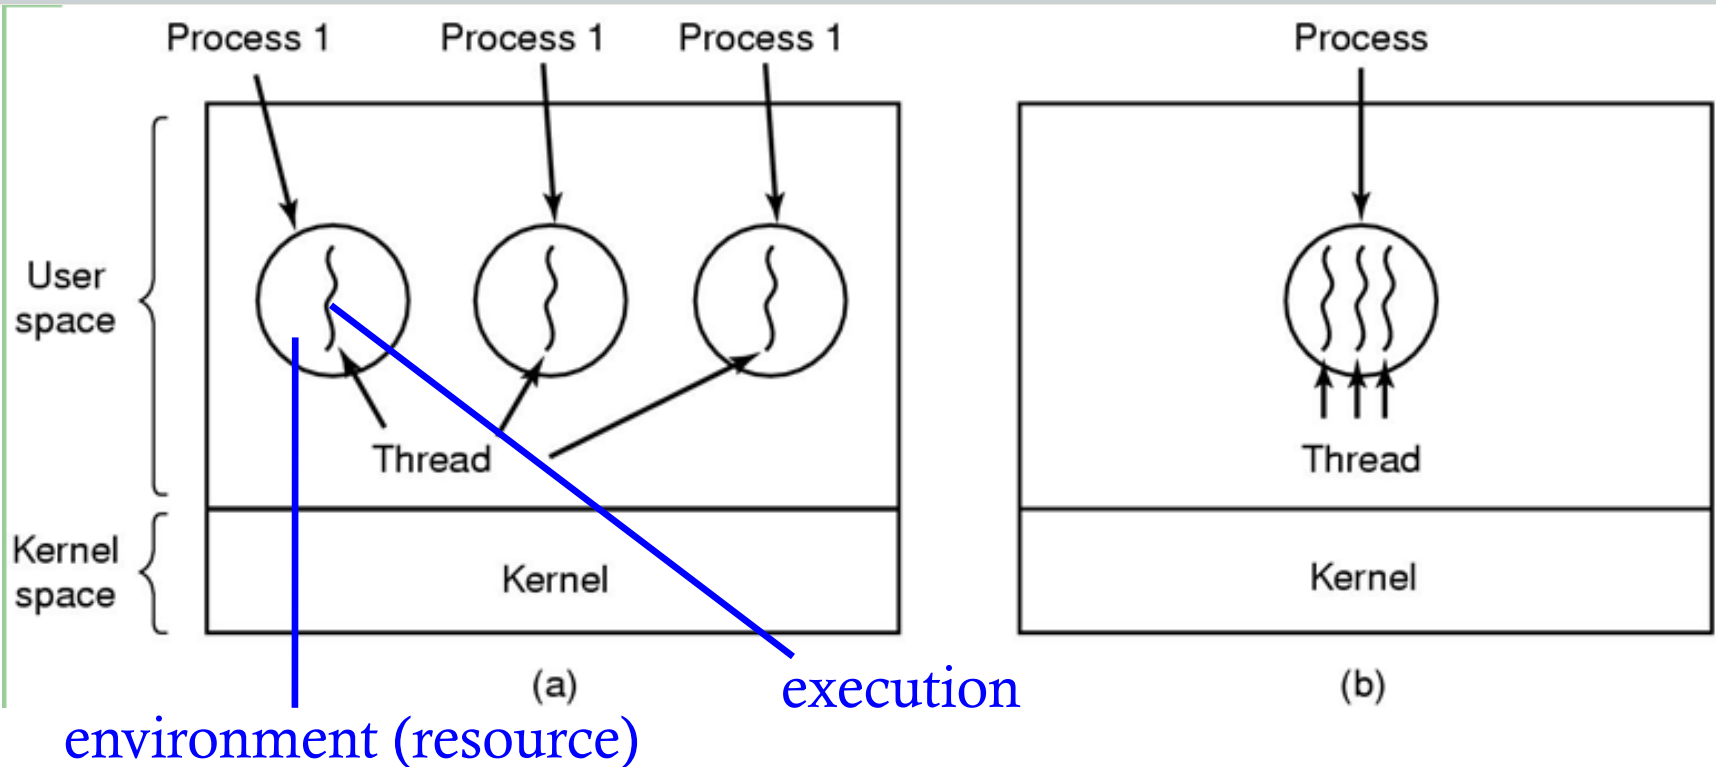
# Rethinking Processes

- What is similar in these cooperating processes?
  - They all share the same code and data (address space)
  - They all share the same privileges
  - They all share the same resources (files, sockets, etc.)
- What don't they share?
  - Each has its own execution state: PC, SP, and registers
- **Key idea:** Why don't we separate the concept of a process from its execution state?
  - **Process:** address space, privileges, resources, etc.
  - **Execution state:** PC, SP, registers
- Exec state also called **thread of control**, or **thread**

# Threads

- Modern OSes (Mac, Windows, modern Unix) separate the concepts of processes and threads
  - The **thread** defines a sequential execution stream within a process (PC, SP, registers)
  - The **process** defines the address space and general process attributes (everything but threads of execution)
- A thread is bound to a single process
  - Processes, however, can have **multiple** threads
- Threads become the unit of scheduling
  - Processes are now the **containers** in which threads execute

# Threads: lightweight processes



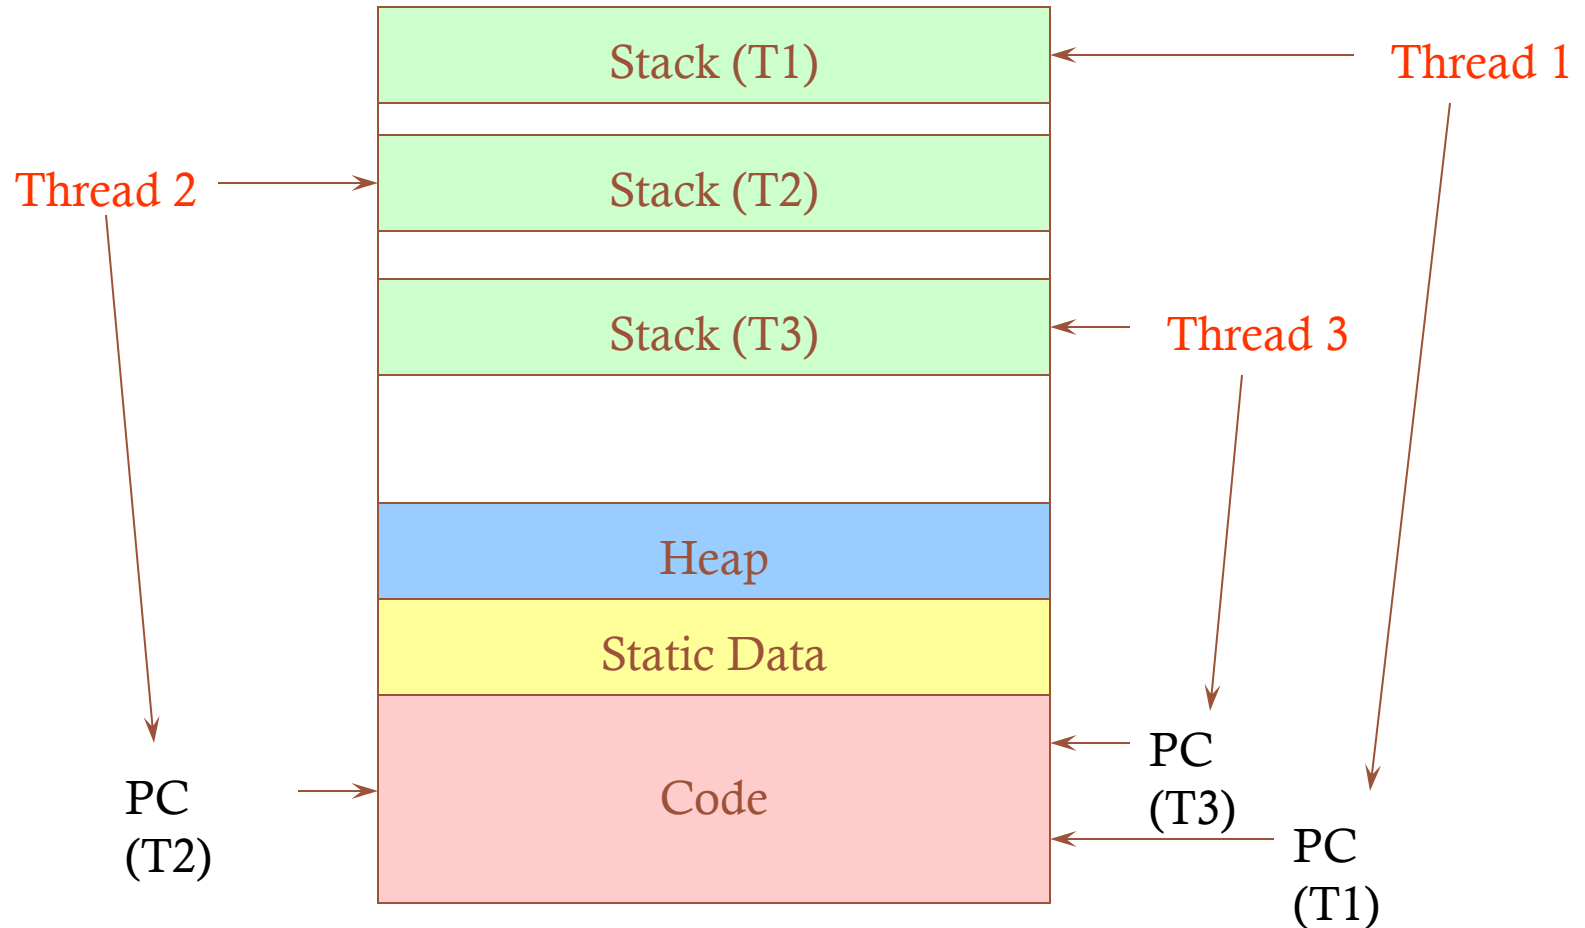
(a) Three processes each with one thread

(b) One process with three threads

# The thread model

- Shared information
  - Processor info: parent process, time, etc
  - Memory: segments, page table, and stats, etc
  - I/O and file: communication ports, directories and file descriptors, etc
- Private state
  - State (ready, running and blocked)
  - Registers
  - Program counter
  - Execution stack
  - Why?
- Each thread execute separately

# Threads in a Process





# Threads: Concurrent Servers

- Using `fork()` to create new processes to handle requests in parallel is overkill for such a simple task
- Recall our forking Web server:

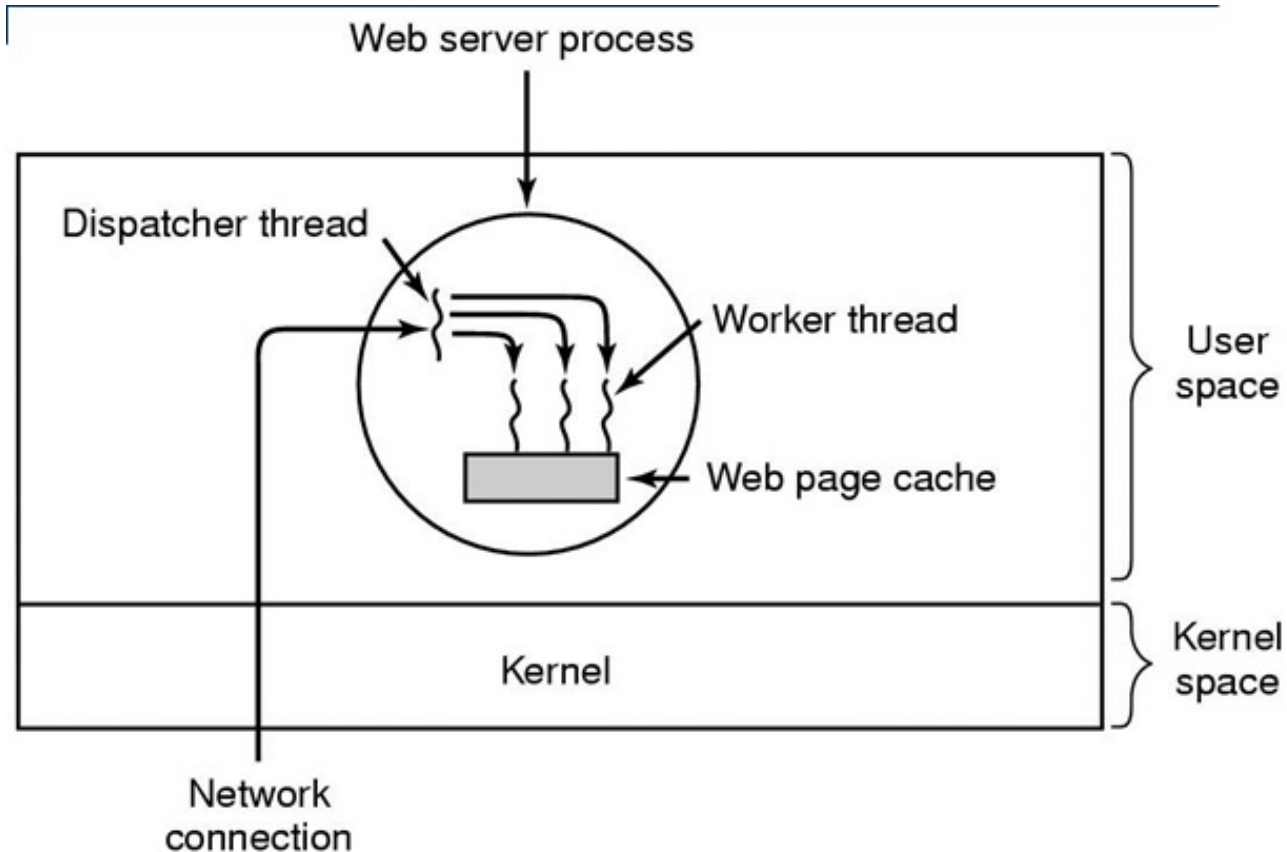
```
while (1) {  
    int sock = accept();  
    if ((child_pid = fork()) == 0) {  
        Handle client request  
        Close socket and exit  
    } else {  
        Close socket  
    }  
}
```

# Threads: Concurrent Servers

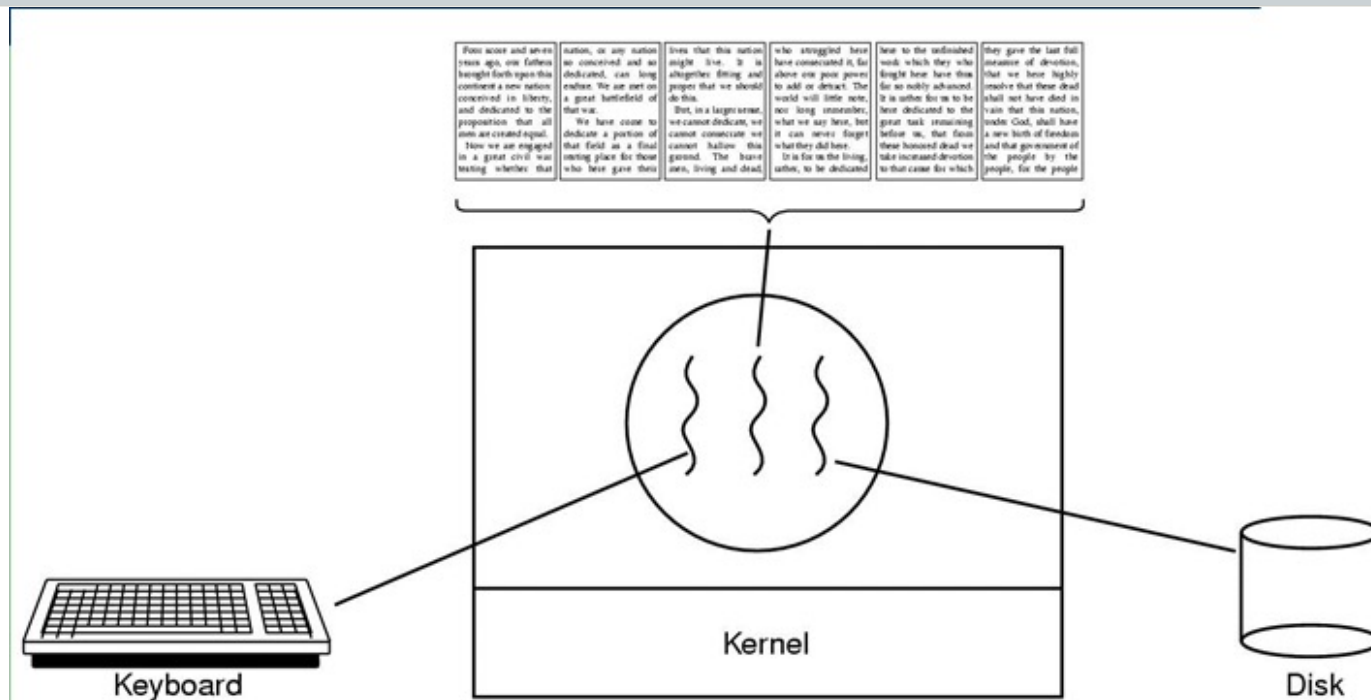
- Instead, we can create a new thread for each request

```
web_server() {  
    while (1) {  
        int sock = accept();  
        thread_create(handle_request, sock);  
    }  
}  
  
handle_request(int sock) {  
    Process request  
    close(sock);  
}
```

# Thread usage: web server



# Thread usage: word processor



- A thread can wait for I/O, while the other threads can still running.
- What if it is single-threaded?

# Kernel-Level Threads

- We have taken the execution aspect of a process and separated it out into threads
  - To make concurrency cheaper
- As such, the OS now manages threads *and* processes
  - All thread operations are implemented in the kernel
  - The OS schedules all of the threads in the system
- OS-managed threads are called **kernel-level threads** or **lightweight processes**
  - Windows: **threads**
  - Solaris: **lightweight processes (LWP)**
  - POSIX Threads (pthreads): **PTHREAD\_SCOPE\_SYSTEM**

# Kernel-level Thread Limitations

- Kernel-level threads make concurrency much cheaper than processes
  - Much less state to allocate and initialize
- However, for fine-grained concurrency, kernel-level threads still suffer from too much overhead
  - Thread operations still require system calls
    - Ideally, want thread operations to be **as fast as a procedure call**
- For such fine-grained concurrency, need even “cheaper” threads

# User-Level Threads

- To make threads cheap and fast, they need to be implemented at user level
  - **Kernel-level threads** are managed by the OS
  - **User-level threads** are managed entirely by the run-time system (user-level library)
- User-level threads are small and fast
  - A thread is simply represented by a PC, registers, stack, and small thread control block (TCB)
  - Creating a new thread, switching between threads, and synchronizing threads are done via **procedure call**
    - No kernel involvement
  - User-level thread operations **100x faster** than kernel threads
  - pthreads: **PTHREAD\_SCOPE\_PROCESS**

# User-level Thread Limitations

- But, user-level threads are not a perfect solution
  - As with everything else, they are a tradeoff
- User-level threads are **invisible** to the OS
  - They are not well integrated with the OS
- As a result, the OS can make poor decisions
  - Scheduling a process with idle threads
  - **Blocking a process whose thread initiated an I/O, even though the process has other threads that can execute**
- Solving this requires communication between the kernel and the user-level thread manager



# Kernel- vs. User-level Threads

- Kernel-level threads
  - Integrated with OS (informed scheduling)
  - Slow to create, manipulate, synchronize
- User-level threads
  - Fast to create, manipulate, synchronize
  - Not integrated with OS (uninformed scheduling)
- Understanding the differences between kernel- and user-level threads is important
  - For programming (correctness, performance)
  - For test-taking

# Kernel- and User-level Threads

- Or use **both** kernel- and user-level threads
  - Can associate a user-level thread with a kernel-level thread
  - Or, multiplex user-level threads on top of kernel-level threads
- Java Virtual Machine (JVM)
  - Java threads used to be user-level threads
  - On older Unix, only one “kernel thread” per process
    - Multiplex all Java threads on this one kernel thread
  - On Windows NT, some more modern Unix
    - Can multiplex Java threads on multiple kernel threads
    - Can have more Java threads than kernel threads
  - No longer the case today
- Golang today uses user-level threads
  - Multiplex multiple Goroutines (user-level threads) on multiple kernel level threads

# Implementing Threads

- Implementing threads has a number of issues
  - Interface
  - Context switch
  - Preemptive vs. Non-preemptive
    - What do they mean?
  - Scheduling
  - Synchronization (next lecture)
- Focus on user-level threads
  - Kernel-level threads are similar to original process management and implementation in the OS

# Sample Thread Interface

- `thread_create`(procedure\_t, arg)
  - Create a new thread of control
  - Start executing procedure\_t
- `thread_yield`()
  - Voluntarily give up the processor
- `thread_exit`()
  - Terminate the calling thread; also thread\_destroy
- `thread_join`(target\_thread)
  - Suspend the execution of calling thread until target\_thread terminates

# Thread Scheduling

- For user-level thread: scheduling occurs entirely in user-space
- The thread scheduler determines when a thread runs
- It uses queues to keep track of what threads are doing
  - Just like the OS and processes
  - But it is implemented at user-level in a library
- Run queue: Threads currently running (usually one)
- Ready queue: Threads ready to run
- Are there wait queues?

# Review of threads

- What are shared among threads of the same process? What are not?
  - Why cannot they share the same stack?
  - How threads of the same process communicate with each other?
- Trade-off between kernel level threads and user level threads?
- Blocking system call
  - Blocking system call: an I/O system call that will wait for the I/O to complete before returning
- How do we implement user-level threads

# Non-Preemptive Scheduling

- Threads voluntarily give up the CPU with `thread_yield`

Ping Thread

```
while (1) {  
    printf("ping\n");  
    thread_yield();  
}
```

Pong Thread

```
while (1) {  
    printf("pong\n");  
    thread_yield();  
}
```

- What is the output of running these two threads?

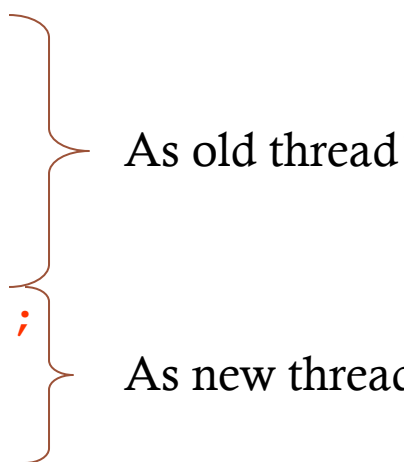
# thread\_yield()

- Wait a second. How does thread\_yield() work?
- The semantics of thread\_yield are that it gives up the CPU to another thread
  - In other words, it **context switches** to another thread
- So what does it mean for thread\_yield to return?
  - It means that *another thread* called thread\_yield!
- Execution trace of ping/pong
  - `printf("ping\n");`
  - `thread_yield();`
  - `printf("pong\n");`
  - `thread_yield();`
  - ...



# Implementing thread\_yield()

```
thread_yield() {  
    thread_t old_thread = current_thread;  
    current_thread = get_next_thread();  
    append_to_queue(ready_queue, old_thread);  
    context_switch(old_thread, current_thread);  
    return;  
}
```



As old thread

As new thread

- The magic step is invoking context\_switch()
- Why do we need to call append\_to\_queue()?

# Thread Context Switch

- The context switch routine does all the magic
  - Saves context of the currently running thread (`old_thread`)
    - Push all machine state onto its stack (*except stack pointer*)
  - Restores context of the next thread
    - Pop all machine state from the next thread's stack
  - The next thread becomes the current thread
  - Return to caller as new thread
- This is all done in assembly language

# Wait a minute

- Non-preemptive threads have to voluntarily give up CPU
  - Only voluntary calls to `thread_yield()`, or `thread_exit()` causes a context switch
- What if a thread never release the CPU (never calls `thread_yield()`)?
- We need preemptive user-level thread scheduling

# Preemptive Scheduling

- **Preemptive scheduling** causes an **involuntary** context switch
  - Need to regain control of processor asynchronously
- How?
  - *Use timer interrupt*
  - Timer interrupt handler forces current thread to “call” `thread_yield`
    - How?

# Process vs. thread

- Multithreading is only an option for “cooperative tasks”
  - Trust and sharing
- Process
  - Strong isolation but poor performance
- Thread
  - Good performance but share too much
- Example: web browsers
  - Safari: multithreading
    - one webpage can crash entire Safari
  - Google Chrome: each tab has its own process

# Summary

- The operating system as a large multithreaded program
  - Each process executes as a thread within the OS
- Multithreading is also very useful for applications
  - Efficient multithreading requires fast primitives
  - Processes are too heavyweight
- Solution is to separate threads from processes
  - Kernel-level threads much better, but still significant overhead
  - User-level threads even better, but not well integrated with OS
- Now, how do we get our threads to correctly cooperate with each other?
  - Synchronization...