

# Operating Systems

## CSCI 3150

### *Lecture 6: Synchronization (II) -- Semaphores and Monitors*

Hong Xu


<https://github.com/henryhxu/CSCI3150>

# Higher-Level Synchronization

- We looked at using locks to provide mutual exclusion
- Locks work, but they have some drawbacks when critical regions are long
  - Spinlocks – inefficient
  - Disabling interrupts – can miss or delay important events
- Instead, we want synchronization mechanisms that
  - Block waiters
  - Leave interrupts enabled inside the critical section
- Look at two common high-level mechanisms
  - **Semaphores**: binary (mutex) and counting
  - **Monitors**: mutexes and condition variables
- Use them to solve common synchronization problems

# Semaphores

- Semaphores are an **abstract data type** that provide mutual exclusion to critical region
- Semaphores can also be used as atomic counters
  - More later
- Semaphores are **integers** that support two operations:
  - wait(semaphore): decrement, block until semaphore is open
    - Also **P**(), after the Dutch word for test, or down()
  - signal(semaphore): increment, allow another thread to enter
    - Also **V**() after the Dutch word for increment, or up()
  - That's it! No other operations – not even just reading its value – exist
- **P** and **V** are probably the most unintuitive names you encounter in this course
  - *and you have Edsger W. Dijkstra to thank to LOL*
- Semaphore safety property: the semaphore value is always greater than or equal to 0



Dijkstra

# Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes/threads
- When **P()** is called by a thread:
  - If semaphore is open ( $> 0$ ), thread continues
  - If semaphore is closed, thread blocks on queue
- Then **V()** opens the semaphore:
  - If a thread is waiting on the queue, the thread is unblocked
    - *What if multiple threads are waiting on the queue?*
  - If no threads are waiting on the queue, the signal is remembered for the next thread
    - In other words, **V()** has “history” (c.f., condition vars later)
    - This “history” is a counter

# Semaphores in OS161

```
P(sem) {  
    Disable interrupts;  
    while (sem->count == 0) {  
        thread_sleep(sem); /* current thread  
                           will sleep on this sem */  
    }  
    sem->count--;  
    Enable interrupts;  
}
```

```
V(sem) {  
    Disable interrupts;  
    sem->count++;  
    thread_wakeup(sem); /* this will wake  
                        up all the threads waiting on this  
                        sem. Why wake up all threads? */  
    Enable interrupts;  
}
```

- thread\_sleep() assumes interrupts are disabled
  - Note that interrupts are disabled only to enter/leave critical section
  - How can it sleep with interrupts disabled?
- What happens if “while (sem->count ==0)” is an “if (sem->count != 0)”?

# Semaphore Types

- Semaphores come in two types
- **Mutex** semaphore (or **binary** semaphore)
  - Represents single access to a resource
  - Guarantees mutual exclusion to a critical section
- **Counting** semaphore (or **general** semaphore)
  - Represents a resource with many units available, or a resource that allows certain kinds of unsynchronized concurrent access (e.g., reading)
  - **Multiple threads can pass the semaphore (P)**
  - Number of threads determined by the semaphore “count”
    - mutex has count = 1, counting has count = N

# Using Semaphores

- Use is similar to our locks, but semantics are different

```
struct Semaphore {  
    int value;  
    Queue q;  
} S;  
  
withdraw (account, amount) {  
    P(S);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    V(S);  
    return balance;  
}
```

**Threads  
block  
critical  
section**

**It is undefined which thread  
runs after a signal**

```
P(S);  
balance = get_balance(account);  
balance = balance - amount;
```

```
P(S);
```

```
P(S);
```

```
put_balance(account, balance);  
V(S);
```

```
...  
V(S);
```

```
...  
V(S);
```



# Possible Deadlocks with Semaphores

Example:

Thread 1:

Thread 2:

share two mutex semaphores S and Q  
 $S := 1; Q := 1;$

P(S);

P(Q);

...

V(Q);

V(S);

P(Q);

P(S);

...

V(S);

V(Q);

Deadlock?





# Semaphore Summary

- Semaphores can be used to solve any of the traditional synchronization problems
- However, they have some drawbacks
  - They are essentially shared global variables
    - Can potentially be accessed anywhere in program
  - No connection between the semaphore and the data being controlled by the semaphore
  - No control or guarantee of proper usage
- Sometimes hard to use and prone to bugs
  - Another approach: Use programming language support

# Monitors

- A monitor is a programming language construct that controls access to shared data
  - Synchronization code added by compiler, enforced at runtime
  - Why is this an advantage?
- A monitor is a module that encapsulates
  - Shared data structures
  - Procedures that operate on the shared data structures
  - Synchronization between concurrent threads that invoke the procedures
- A monitor protects its data from unstructured access
- It guarantees that threads accessing its data through its procedures interact only in legitimate ways

# Monitor Semantics

- A monitor guarantees mutual exclusion
  - Only one thread can execute any monitor procedure at any time (the thread is “in the monitor”)
  - If a second thread invokes a monitor procedure when a first thread is already executing one, it blocks
    - So the monitor has to have a wait queue...
  - If a thread within a monitor blocks, another one can enter
    - Condition Variable
- What are the implications in terms of parallelism in monitor?

# Account Example

```
Monitor account {  
    double balance;  
  
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }  
}
```

**Threads  
block  
waiting  
to get  
into  
monitor**

**When first thread exits, another can  
enter. Which one is undefined.**

**withdraw(amount)**  
balance = balance - amount;

withdraw(amount)

withdraw(amount)

return balance (and exit)

balance = balance - amount  
return balance;

balance = balance - amount;  
return balance;

- Hey, that was easy
- But what if a thread wants to wait inside the monitor?

# Condition Variables

- A **condition variable** is associated with a condition needed for a thread to make progress once it is in the monitor.

```
Monitor M {  
  ... monitored variables  
  Condition c;
```

```
void enter_mon (...) {  
  if (extra property not true) wait(c);  
  do what you have to do  
  if (extra property true) signal(c);  
}
```

waits outside of the monitor's mutex

brings in one thread waiting on condition

# Condition Variables

- Condition variables support three operations:
  - **Wait** – release monitor lock, wait for C/V to be signaled
    - So condition variables have wait queues, too
  - **Signal** – wakeup one waiting thread
  - **Broadcast** – wakeup all waiting threads
- Condition variables *are not* boolean objects
  - “if (condition\_variable) then” ... does not make sense
  - “if (num\_resources == 0) then wait(resources\_available)” does
  - An example will make this more clear

# Condition Vars != Semaphores

- Condition variables != semaphores
  - However, they each can be used to implement the other
- Access to the monitor is controlled by a lock
  - wait() blocks the calling thread, and gives up the lock
    - To call wait, the thread has to be in the monitor (hence has lock)
    - Semaphore::P just blocks the thread on the queue
  - signal() causes a waiting thread to wake up
    - If there is no waiting thread, the signal is lost
    - Semaphore::V increases the semaphore count, allowing future entry even if no thread is waiting
    - Condition variables have no history



# Locks and Condition Vars

- A C.V. allows a thread to be signaled when some condition is satisfied. By itself, mutex (and lock) doesn't do this.
- A queue of items to work on
  - Use a lock to ensure mutex to access the queue
  - Need C.V. to tell a consumer thread that the queue is non-empty or empty
    - Otherwise, threads have to poll the queue...

# Using Semaphores

- We've looked at a simple example for using synchronization
  - Mutual exclusion while accessing a bank account
- Now we're going to use semaphores to look at more interesting examples
  - Readers/Writers
  - Bounded Buffers (after we discuss Monitor)

# Readers/Writers Problem

- Readers/Writers Problem:
  - An object is shared among several threads
  - Some threads only read the object, others only write it
  - We can allow **multiple readers** but only **one writer**
    - Let  $\#r$  be the number of readers,  $\#w$  be the number of writers
    - Safety:  $(\#r \geq 0) \wedge (0 \leq \#w \leq 1) \wedge ((\#r > 0) \Rightarrow (\#w = 0))$
- How can we use semaphores to control access to the object to implement this protocol?

# First write operational code

```
reader {  
  read;  
}
```

```
writer {  
  Write;  
}
```

- Does it work?
- Why?

# First attempt: one mutex semaphore

```
// exclusive writer or reader
Semaphore w_or_r = 1;

reader {
    P(w_or_r); // lock out writers
    read;
    V(w_or_r); // up for grabs
}

writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

- Does it work?
- Why?
- Which condition is satisfied and which is not?  
( $\#r \geq 0$ )  
( $0 \leq \#w \leq 1$ )  
( $(\#r > 0) \Rightarrow (\#w = 0)$ )

# Second attempt: add a counter

```
int readcount = 0; // record #readers
Semaphore w_or_r = 1; // mutex semaphore
```

```
reader {
    readcount++;
    if (readcount == 1){
        P(w_or_r); // lock out writers
    }
    read;
    readcount--;
    if (readcount == 0){
        V(w_or_r); // up for grabs
    }
}
```

```
writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

- Does it work?

- *readcount is a shared variable, who protects it?*

**Thread 1:**

```
reader {
    readcount++;
```

context switch

```
if (readcount == 1){
    P(w_or_r);
}
```

**Thread 2:**

```
reader {
    readcount++;
    if (readcount == 1){
        P(w_or_r);
    }
```

A context switch can happen, a writer can come in since no reader locked the semaphore!

# Readers/Writers Real Solution

- Use three variables
  - int `readcount` – number of threads reading object
  - Semaphore `mutex` – control access to readcount
  - Semaphore `w_or_r` – exclusive writing or reading



# Readers/Writers

```
// number of readers
int readcount = 0;
// mutual exclusion to readcount
Semaphore mutex = 1;
// exclusive writer or reader
Semaphore w_or_r = 1;

writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

```
reader {
    P(mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        P(w_or_r); // synch w/ writers
    V(mutex); // unlock readcount
    Read;
    P(mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        V(w_or_r); // up for grabs
    V(mutex); // unlock readcount
}
```

- Why do readers use **mutex**?
- What if the **V(mutex)** is above “if (readcount == 1)”?

# But it still has a problem...

```
// number of readers
int readcount = 0;
// mutual exclusion to readcount
Semaphore mutex = 1;
// exclusive writer or reader
Semaphore w_or_r = 1;

writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

```
reader {
    P(mutex);    // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        P(w_or_r); // synch w/ writers
    V(mutex);    // unlock readcount
    Read;
    P(mutex);    // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        V(w_or_r); // up for grabs
    V(mutex);    // unlock readcount
}
```

# Problem: Starvation

- What if a writer is waiting, but readers keep coming, the writer is starved



# Semaphore Questions

- Are there any problems that **can be solved** with counting semaphores that **cannot be solved** with mutex semaphores?
- If a system provides only mutex semaphores, can you use it to implement a counting semaphores?
- When to use counting semaphore?
  - Problem needs a counter
  - The maximum value is known (bounded)

# Monitor Readers and Writers

- Will have four methods: **StartRead**, **StartWrite**, **EndRead** and **EndWrite**
- Monitored data: **nr** (number of readers) and **nw** (number of writers) with the monitor invariant

$$(nr \geq 0) \wedge (0 \leq nw \leq 1) \wedge ((nr > 0) \Rightarrow (nw = 0))$$

- Two conditions:
  - **canRead**:  $nw = 0$
  - **canWrite**:  $(nr = 0) \wedge (nw = 0)$

# Monitor Readers and Writers

```
Monitor RW {  
    int nr = 0, nw = 0;  
    Condition canRead, canWrite;  
  
    void StartRead () {  
        while (nw != 0) do wait(canRead);  
        nr++;  
    }  
  
    void EndRead () {  
        nr--;  
        if (nr == 0) signal(canWrite);  
    }  
}
```

```
void StartWrite {  
    while (nr != 0 || nw != 0) do wait(canWrite);  
    nw++;  
}  
  
void EndWrite () {  
    nw--;  
    broadcast(canRead);  
    signal(canWrite);  
}  
} // end monitor
```

# Monitor Readers and Writers

- Is there any priority between readers and writers?
- What if you wanted to ensure that a waiting writer would have priority over new readers?



# Bounded Buffer

- Problem: There is a set of resource buffers shared by producer and consumer threads
  - **Producer** inserts resources into the buffer set
    - Output, disk blocks, memory pages, processes, etc.
  - **Consumer** removes resources from the buffer set
    - Whatever is generated by the producer
- Producer and consumer execute at different rates
  - No serialization of one behind the other
  - Tasks are independent (easier to think about)
  - The buffer set allows each to run without explicit handoff
- Safety:
  - Sequence of consumed values is prefix of sequence of produced values
  - If  $nc$  is number consumed,  $np$  number produced, and  $N$  the size of the buffer, then  $0 \leq np - nc \leq N$

# Bounded Buffer (2) – functional code

```
producer {  
  while (1) {  
    Produce new resource;  
    Add resource to an empty buffer;  
  }  
}
```

```
consumer {  
  while (1) {  
    Remove resource from a full buffer;  
    Consume resource;  
  }  
}
```

# Bounded Buffer (3)

- Use three semaphores:
  - **empty** – count of empty buffers
    - Counting semaphore
    - $\text{empty} = N - (np - nc)$
  - **full** – count of full buffers
    - Counting semaphore
    - $np - nc = \text{full}$
  - **mutex** – mutual exclusion to shared set of buffers
    - Binary semaphore

# Bounded Buffer (4)

```
Semaphore mutex = 1; // mutual exclusion to shared set of buffers
Semaphore empty = N; // count of empty buffers (all empty to start)
Semaphore full = 0;   // count of full buffers (none full to start)
```

```
producer {
    while (1) {
        Produce new resource;
        P(empty); // wait for empty buffer
        P(mutex); // lock buffer list
        Add resource to an empty buffer;
        V(mutex); // unlock buffer list
        V(full);  // note a full buffer
    }
}
```

```
consumer {
    while (1) {
        P(full); // wait for a full buffer
        P(mutex); // lock buffer list
        Remove resource from a full buffer;
        V(mutex); // unlock buffer list
        V(empty); // note an empty buffer
        Consume resource;
    }
}
```

# Bounded Buffer (5)

Producer



Producer decrements EMPTY and blocks when buffer is full since the semaphore is at 0

Consumer



Consumer decrements FULL and blocks when buffer has no item since the semaphore FULL is at 0

# Bounded Buffer (6)

*Why we need both “empty” and “full” semaphores?*

```
Semaphore mutex = 1; // mutual exclusion to shared set of buffers
Semaphore empty = N; // count of empty buffers (all empty to start)
Semaphore full = 0; // count of full buffers (none full to start)
```

```
producer {
  while (1) {
    Produce new resource;
    P(empty); // wait for empty buffer
    P(mutex); // lock buffer list
    Add resource to an empty buffer;
    V(mutex); // unlock buffer list
    V(full); // note a full buffer
  }
}
```

```
consumer {
  while (1) {
    P(full); // wait for a full buffer
    P(mutex); // lock buffer list
    Remove resource from a full buffer;
    V(mutex); // unlock buffer list
    V(empty); // note an empty buffer
    Consume resource;
  }
}
```

*More consumers “remove resource” than actually produced!*

# Monitor Bounded Buffer

```
Monitor bounded_buffer {  
    Resource buffer[N];  
    // Variables for indexing buffer  
    // monitor invariant involves these vars  
    Condition not_full; // space in buffer  
    Condition not_empty; // value in buffer  
  
    void put_resource (Resource R) {  
        while (buffer array is full)  
            wait(not_full);  
        Add R to buffer array;  
        signal(not_empty);  
    }  
}
```

```
Resource get_resource() {  
    while (buffer array is empty)  
        wait(not_empty);  
    Get resource R from buffer array;  
    signal(not_full);  
    return R;  
}  
} // end monitor
```

- What happens if no threads are waiting when signal is called?
  - Signal is lost



# Monitor Queues

```
Monitor bounded_buffer {
```

```
    Condition not_full;
```

```
    ...other variables...
```

```
    Condition not_empty;
```

```
    void put_resource () {
```

```
        ...wait(not_full)...
```

```
        ...signal(not_empty)...
```


```
    }
```

```
    Resource get_resource () {
```

```
        ...
```

```
    }
```

```
}
```



# Summary

- Semaphores
  - P()/V() implement blocking mutual exclusion
  - Also used as atomic counters (counting semaphores)
  - Can be inconvenient to use
- Monitors
  - Synchronizes execution within procedures that manipulate encapsulated data shared among procedures
    - Only one thread can execute within a monitor at a time
  - Relies upon high-level language support
- Condition variables
  - Used by threads as a synchronization point to wait for events
  - Inside monitors

<https://github.com/henryhxu/CSCI3150>